

# 深度强化学习

Week 14 Deep Q-learning Network

# 实验任务

- 用Deep Q-learning Network(DQN)玩CartPole-v1游戏， 框架代码已经给出，至少需要补充‘TODO’标记的代码片段。

- 结果要求与展示


- 至少完成：
  - 500局(Episodes)游戏内，达成一次：**连续10局Reward值为500**
  - 展示：“单局Reward值”曲线以及“最近100局的平均Reward值”曲线
- 进阶：
  - 达成一次：“**最近百局的平均Reward值**”  $\geq 475$
  - 更快地达到这个目标
  - 更高的百局平均Reward值
  - ...

- Deadline

- 三周， **6月17日23:59**

- 提交格式

- E8\_学号.zip



# 续

- 需要补充的代码包括
  - Qnet
    - 补充一个线性层
  - ReplayBuffer
    - 所有成员函数的实现
  - DQN
    - choose\_action
      - $\epsilon$ -greedy策略代码
    - learn
      - Q值的计算
      - 目标值计算
      - 损失值计算
      - 梯度下降
- 可根据需要进一步调整/改进

## Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

With probability  $\epsilon$  select a random action  $a_t$   
otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$



Every  $C$  steps reset  $\hat{Q} = Q$

**End For**

**End For**

# CartPole-v1

- 创建环境(gym==0.23)
  - `env = gym.make("CartPole-v1")`
- 重置环境， 并获取初始状态
  - `obs = env.reset()`
- 执行动作， 并获取下一步状态， 当前奖励， 是否结束等信息
  - `next_obs, reward, done, info = env.step(action)`
- `obs`
  - [车位置， 车速度， 杆角度， 杆角速度]
- `action`
  - 离散值0/1， 表示向左/右移动
- `reward`
  - 不结束时都为1
  - 可根据需要进行修改（奖励工程）



# Deep Q-learning Network(DQN)

- 用网络预测Q值，即网络输入状态state，输出该状态下每个动作的Q值
- Q值的更新变为网络参数的更新，因此网络的损失值可定义为均方误差

$$L = \frac{1}{2} \left( Q(s, a) - \left( r + \gamma \max_{a'} Q(s', a') \right) \right)^2$$

- 探索与利用
  - $\epsilon - greedy$  以概率 $\epsilon$ 随机选择一个动作，以概率 $1 - \epsilon$ 选择最佳动作
- 经验回放
  - 用回放缓存区可以减少与环境做互动的次数，提高训练效率
  - 减少同批次训练数据的依赖关系
  - 做法：将每一步转移的状态、动作、奖励、下一状态等信息存到一个缓冲区，训练网络时从缓冲区随机抽取一批数据作为训练数据

---

**Algorithm 1** Deep Q-learning with Experience Replay

---

```
Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights
for episode = 1,  $M$  do
  Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$ 
  for  $t = 1, T$  do
    With probability  $\epsilon$  select a random action  $a_t$ 
    otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$ 
    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ 
    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ 
    Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$ 
    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3
  end for
end for
```

---

[DQN2013版](#)（无Target网络）

# Deep Q-learning Network(DQN)

- 网络的损失值

- $$L = \frac{1}{2} \left( Q(s, a) - \left( r + \gamma \max_{a'} Q(s', a') \right) \right)^2$$

- 目标网络

- $\left( r + \gamma \max_{a'} Q(s', a') \right)$ 可看作目标值，目标值跟随Q一直变化会给训练带来困难
- 将评估网络与目标网络分开，目标网络不训练，评估网络每更新若干轮后，用评估网络参数替换目标网络参数

- $$L = \frac{1}{2} \left( Q_{eval}(s, a) - \left( r + \gamma \max_{a'} Q_{target}(s', a') \right) \right)^2$$

## Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

        With probability  $\varepsilon$  select a random action  $a_t$

        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

**End For**

**End For**

[DQN2015版](#) (有Target网络)